

# Customer Churn Case Study

We will take the case study through 3 steps in order to compute a good fit model for predicting a customer churn  
Analyzed for candidacy - Senior Manager - Business Intelligence at ZoomInfo - Powered by DiscoverOrg



1. Data Summary and Feature Selection Using three methods
2. Separating out Train and Test Data
3. Applying model to Train Data and test resulting model on Test Data 80/20 split respectively

```
In [60]: #Importing required packages to run a Univariate Feature Selection method
import pandas as pd
import numpy as np

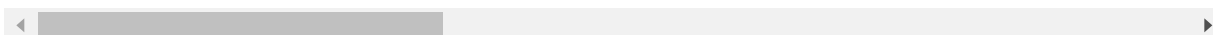
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
from sklearn.ensemble import ExtraTreesClassifier
```

```
In [61]: #Loading data
churn_data = pd.read_csv(r"C:\Users\lemic\Downloads\telco-customer-churn\WA_Fn-UseC_-Telco-Customer-Churn.csv")
churn_data.head()
```

Out[61]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service
1	5575-GNVDE	Male	0	No	No	34	Yes	No
2	3668-QPYBK	Male	0	No	No	2	Yes	No
3	7795-CFOCW	Male	0	No	No	45	No	No phone service
4	9237-HQITU	Female	0	No	No	2	Yes	No

5 rows × 21 columns



```
In [16]: #print info on dataset
churn_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
customerID      7043 non-null object
gender          7043 non-null object
SeniorCitizen   7043 non-null int64
Partner         7043 non-null object
Dependents      7043 non-null object
tenure          7043 non-null int64
PhoneService    7043 non-null object
MultipleLines   7043 non-null object
InternetService 7043 non-null object
OnlineSecurity  7043 non-null object
OnlineBackup    7043 non-null object
DeviceProtection 7043 non-null object
TechSupport     7043 non-null object
StreamingTV     7043 non-null object
StreamingMovies 7043 non-null object
Contract        7043 non-null object
PaperlessBilling 7043 non-null object
PaymentMethod   7043 non-null object
MonthlyCharges  7043 non-null float64
TotalCharges    7043 non-null float64
Churn           7043 non-null object
dtypes: float64(2), int64(2), object(17)
memory usage: 1.1+ MB
```

Based on info output, looks like total charges, being a float64 type has either some nulls or non integers values not being recognized

```
In [20]: churn_data = churn_data.drop('customerID', axis=1) #not needed for analysis
churn_data['TotalCharges'] = pd.to_numeric(churn_data['TotalCharges'], errors
= 'coerce') #convert to numeric where possible - got an error for " " nulls so
will fill with 0
churn_data[churn_data['TotalCharges'].isna() == True] = 0 # setting nas to zero
based on finding running to_numeric error
```

```

In [21]: #applyign SelectKBest with chi2 fmethod to select best fit features
import matplotlib.pyplot as plt

#extract subset of object data types
obj_churn_data = churn_data.select_dtypes(include=['object']).copy()

# obj_churn_data[obj_churn_data.isnull().any(axis=1)] #no nulls, thats good!

#obj_churn_data.describe()
obj_churn_data_numeric = pd.get_dummies(obj_churn_data)
#obj_churn_data_numeric.info()

int_churn_data = churn_data.select_dtypes(exclude=['object']).copy()

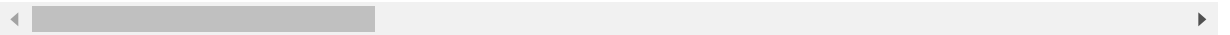
feature_churn = pd.concat([int_churn_data,obj_churn_data_numeric], axis=1)
feature_churn.head()

```

Out[21]:

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	gender_0	gender_Female	gender_Male
0	0	1	29.85	29.85	0	1	0
1	0	34	56.95	1889.50	0	0	1
2	0	2	53.85	108.15	0	0	1
3	0	45	42.30	1840.75	0	0	1
4	0	2	70.70	151.65	0	1	0

5 rows × 63 columns



```
In [263]: #setting variables for feature selection
X = feature_churn.iloc[:,0:-2] #independent columns for univariate selection
Y = feature_churn.iloc[:, -1] #target column which is last column

BestUnivariateFeatures = SelectKBest(score_func=chi2, k=10)
fit = BestUnivariateFeatures.fit(X,Y)
dfscores = pd.DataFrame(fit.scores_)
dfcolumns = pd.DataFrame(X.columns)
featurescores = pd.concat([dfcolumns,dfscores],axis=1)
featurescores.columns = ['feature-univariate-chi2','score']
top_10_univariate_feats = featurescores.nlargest(10, columns = 'score')
top_10_univariate_feats['rank'] = top_10_univariate_feats['score'].rank(ascending=False)
print(top_10_univariate_feats)
```

	feature-univariate-chi2	score	rank
3	TotalCharges	624292.003004	1.0
1	tenure	16278.923685	2.0
2	MonthlyCharges	3733.878622	3.0
49	Contract_Month-to-month	519.895311	4.0
51	Contract_Two year	485.007178	5.0
58	PaymentMethod_Electronic check	426.422767	6.0
25	OnlineSecurity_No	416.716774	7.0
37	TechSupport_No	406.645334	8.0
22	InternetService_Fiber optic	374.476216	9.0
29	OnlineBackup_No	284.531747	10.0

We can see that the top three most prominent features that help determine customer churn (value = 1) are Total charges, Tenure, and Monthly Charges

The trailing features that could play an important part in the final stretch of the predictive model are whether they have: 49 Contract\_Month-to-month 519.895311 4.0 51 Contract\_Two year 485.007178 5.0 58 PaymentMethod\_Electronic check 426.422767 6.0 25 OnlineSecurity\_No 416.716774 7.0 37 TechSupport\_No 406.645334 8.0 22 InternetService\_Fiber optic 374.476216 9.0 29 OnlineBackup\_No 284.531747 10.0

```

In [68]: #feature detection/selection using extratreesclassifier method
model_churn = ExtraTreesClassifier(n_estimators = 100) #100 trees in forest
model_churn.fit(X,Y)
#print(model_churn.feature_importances_) #inbuilt class from sklearn
feat_importance = pd.Series(model_churn.feature_importances_,index=X.columns)
feat_importance = pd.DataFrame(feat_importance)
feat_importance.reset_index(inplace = True)
feat_importance.columns = ['feature-extra-trees','Classifier Importance']
feat_tree_top10 = feat_importance.nlargest(10, columns = 'Classifier Importance')
plt.barh(feat_tree_top10['feature-extra-trees'], feat_tree_top10['Classifier Importance'],align='edge')
plt.gca().invert_yaxis()
plt.show

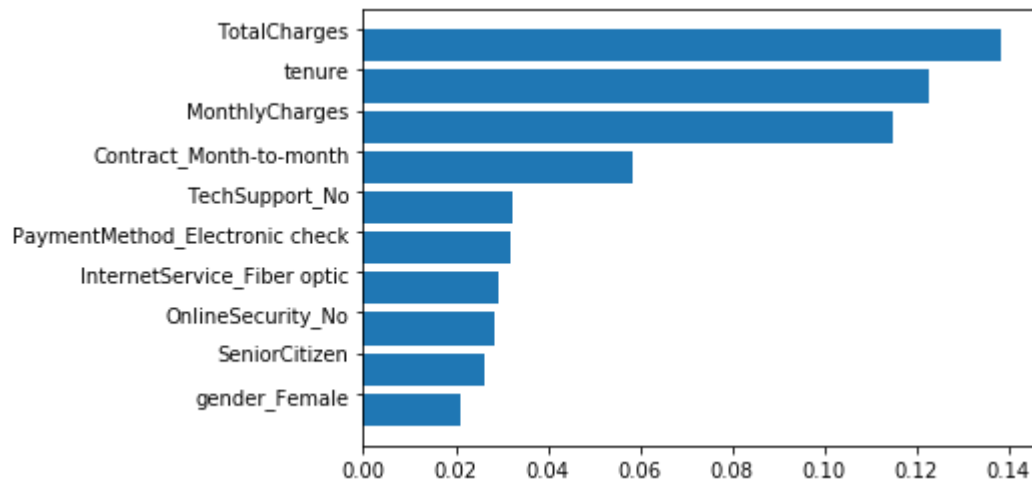
#curious about less important features? print below by uncommenting and running
#print(feat_importance.nsmallest(40, columns = 'Classifier Importance'))

```

```

Out[68]: <function matplotlib.pyplot.show(*args, **kw)>

```



We now have the top 12 important features using the ensemble extra trees classifiers within a forest with 100 trees

In [69]: *#compare of univariate vs ensemble method*

```
pd.concat([top_10_univariate_feats, feat_tree_top10], axis = 1).sort_values(by
='Classifier Importance', ascending=False) # allows comparison between univariate
feature selection vs extratreesclassifier feature detection
```

Out[69]:

	feature-univariate-chi2	score	rank	feature-extra-trees	Classifier Importance
3	TotalCharges	624292.003004	1.0	TotalCharges	0.138462
1	tenure	16278.923685	2.0	tenure	0.122828
2	MonthlyCharges	3733.878622	3.0	MonthlyCharges	0.114840
49	Contract_Month-to-month	519.895311	4.0	Contract_Month-to-month	0.058625
37	TechSupport_No	406.645334	8.0	TechSupport_No	0.032556
58	PaymentMethod_Electronic check	426.422767	6.0	PaymentMethod_Electronic check	0.031774
22	InternetService_Fiber optic	374.476216	9.0	InternetService_Fiber optic	0.029503
25	OnlineSecurity_No	416.716774	7.0	OnlineSecurity_No	0.028391
0	NaN	NaN	NaN	SeniorCitizen	0.026131
5	NaN	NaN	NaN	gender_Female	0.021046
29	OnlineBackup_No	284.531747	10.0	NaN	NaN
51	Contract_Two year	485.007178	5.0	NaN	NaN

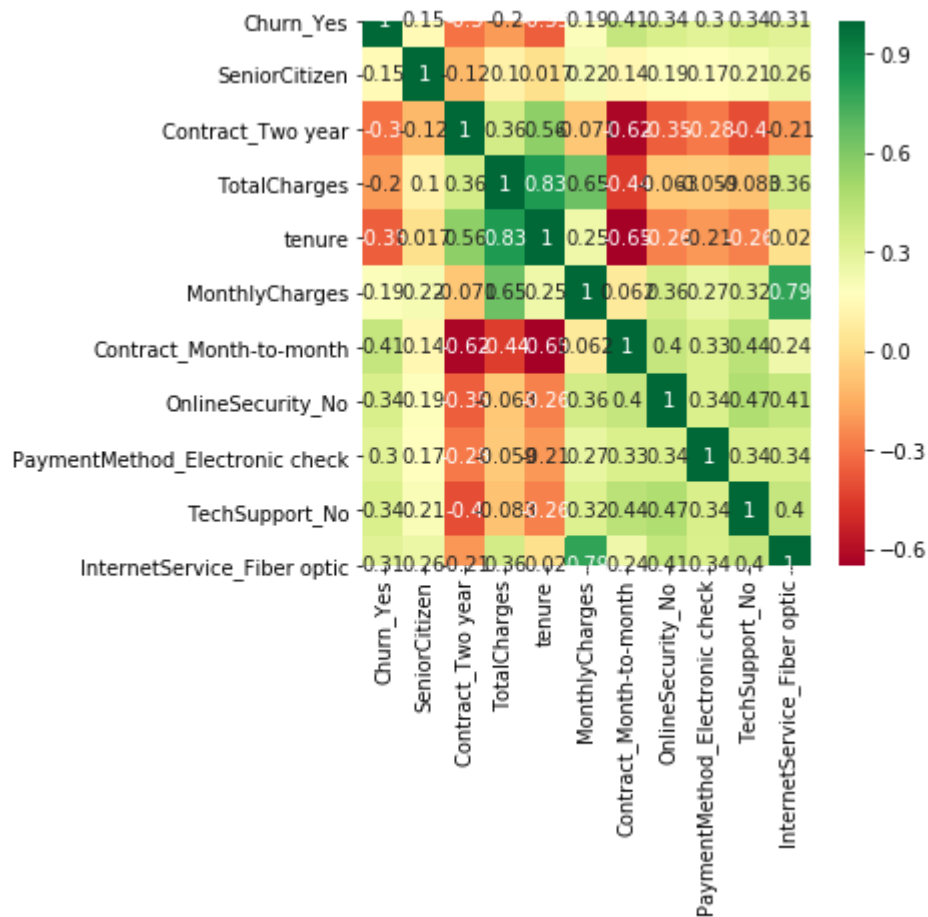
Based on findings above, I will train a model based on the top matching features, which are:

1. TotalCharges 624292.003004 1.0 TotalCharges 0.136892
2. tenure 16278.923685 2.0 tenure 0.123686
3. MonthlyCharges 3733.878622 3.0 MonthlyCharges 0.115721
4. Contract\_Month-to-month 519.895311 4.0 Contract\_Month-to-month 0.060624
5. OnlineSecurity\_No 416.716774 7.0 OnlineSecurity\_No 0.031221
6. PaymentMethod\_Electronic check 426.422767 6.0 PaymentMethod\_Electronic check 0.029601
7. TechSupport\_No 406.645334 8.0 TechSupport\_No 0.028722
8. InternetService\_Fiber optic 374.476216 9.0 InternetService\_Fiber optic 0.026591

In [234]: *#As a treat, let's show a correlation heat map to see how those features affect the churn rate*  
**import seaborn as sns**

In [261]: `selected_feats_churn = feature_churn[['Churn_Yes', 'SeniorCitizen', 'Contract_Two year', 'TotalCharges', 'tenure', 'MonthlyCharges', 'Contract_Month-to-month', 'OnlineSecurity_No', 'PaymentMethod_Electronic check', 'TechSupport_No', 'InternetService_Fiber optic']]`  
`corrmat = selected_feats_churn.corr()`  
`print(len(corrmat))`

```
In [262]: top_corr_features = corrmat.index
g = sns.heatmap(feature_churn[top_corr_features].corr(),annot=True,cmap="RdYlGn")
#top and bottom row cut off due to unstable seaborn and matplotlib with latest
update. could revert to older version of matplotlib, but we'll roll with this
for now.
```



Of the top 10 features identified, 7 have a positive correlation to a customer churning and 3 have a negative correlation

## Summary of Feature Selection

Based on findings, an analyst would find most value building visualizations or reports that target the 8 features identified above, and when there is a deviation above the median for numerical measurements, then a flag or warning should appear next to that customer record. This would include TotalCharges, Tenure, MonthlyCharges. For the categorical variables, having a contract month-to-month and a combination of one or more of the above breaching a median, would flag the customer as a potential churn, and allow sales rep to reach out and attempt a conversion to contract vs month-to-month. The same would go for all other categorical features such as tech support and online security. Improving the online security of a user, and tech support are simple steps that the company can take to retain more customers as shown below on the correlation map. (having online security and tech support reduces churn)

```
Churn_Yes          1.000000
OnlineSecurity_Yes -0.170573
TechSupport_Yes    -0.164016
```

## plot data

```
from pylab import rcParams
import warnings
warnings.filterwarnings("ignore")
```

```
sizes = feature_churn['Churn_Yes'].value_counts(sort=True)
colors = ["grey", "blue"]
rcParams['figure.figsize'] = 5,5
```

## Pie Chart

```
plt.pie(sizes, explode = (0.05,0.05), labels = ["Yes", "No"], colors = colors, autopct = '%1.1f%%', startangle = 180)
```

```
plt.title('Percentage Churn')
plt.show()
```

## Logistic Regression Model to test features (Baseline)

```
### Model shows 80% accuracy in predicting outcomes of Churn based on inputs. Running optimization statistical
logistics regression model with X_feat = selected_feats_churn.loc[:,1:] X_train, X_test, Y_train, Y_test =
train_test_split(X_feat,Y,test_size = 0.30, random_state = 101)
import statsmodels.api as sm
logit_model=sm.Logit(Y_train,X_train)
result=logit_model.fit()
print(result.summary2())
#we get the below summary
Removing Month to Month as it is not statistically significant
Optimization terminated successfully. Current function value: inf Iterations 8 Results: Logit
===== Model:
Logit Pseudo R-squared: inf Dependent Variable: Churn_Yes AIC: inf Date: 2019-11-22 17:52 BIC: inf No.
Observations: 4930 Log-Likelihood: -inf Df Model: 9 LL-Null: 0.0000 Df Residuals: 4920 LLR p-value: 1.0000
Converged: 1.0000 Scale: 1.0000 No. Iterations: 8.0000 -----
Coef. Std.Err. z P>|z| [0.025 0.975] ----- SeniorCitizen 0.2916
0.0982 2.9687 0.0030 0.0991 0.4841 Contract_Two year -1.2530 0.2057 -6.0911 0.0000 -1.6562 -0.8498
TotalCharges 0.0007 0.0001 9.6265 0.0000 0.0006 0.0008 tenure -0.0932 0.0064 -14.5685 0.0000 -0.1057 -0.0806
MonthlyCharges -0.0173 0.0023 -7.3935 0.0000 -0.0219 -0.0127 Contract_Month-to-month -0.0605 0.0964 -0.6272
```



0.5305 -0.2494 0.1285 OnlineSecurity\_No 0.4384 0.0932 4.7031 0.0000 0.2557 0.6211 PaymentMethod\_Electronic  
check 0.4555 0.0812 5.6073 0.0000 0.2963 0.6147 TechSupport\_No 0.4018 0.0929 4.3242 0.0000 0.2197 0.5839  
InternetService\_Fiber optic 1.2088 0.1330 9.0912 0.0000 0.9482 1.4694

```
In [271]: #model with selected features:
feat_lin = feature_churn[['Churn_Yes', 'SeniorCitizen', 'Contract_Two year', 'TotalCharges', 'tenure', 'MonthlyCharges', 'OnlineSecurity_No', 'PaymentMethod_Electronic check', 'TechSupport_No', 'InternetService_Fiber optic']]
X_feat = feat_lin.iloc[:,1:]
X_train, X_test, Y_train, Y_test = train_test_split(X_feat,Y,test_size = 0.30,
random_state = 101)

import statsmodels.api as sm
logit_model=sm.Logit(Y_train,X_train)
result=logit_model.fit()
print(result.summary2())
#we get the below summary

#import model
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
result = model.fit(X_train, Y_train)

from sklearn import metrics
prediction_test = model.predict(X_test)

#print accuracy
print(str(metrics.accuracy_score(Y_test, prediction_test)*100)+'%')
```

Optimization terminated successfully.

Current function value: inf

Iterations 8

Results: Logit

```
=====
==
Model:                Logit                Pseudo R-squared:      inf
Dependent Variable:   Churn_Yes              AIC:                  inf
Date:                2019-11-22 18:11        BIC:                  inf
No. Observations:     4930                  Log-Likelihood:       -inf
Df Model:             8                    LL-Null:              0.00
00
Df Residuals:         4921                  LLR p-value:          1.00
00
Converged:            1.0000                Scale:                1.00
00
No. Iterations:       8.0000
-----
--
                        Coef.  Std.Err.   z      P>|z|   [0.025  0.97
5]
-----
--
SeniorCitizen          0.2883   0.0981   2.9399 0.0033   0.0961   0.48
05
Contract_Two year      -1.2321  0.2032  -6.0640 0.0000  -1.6304  -0.83
39
TotalCharges           0.0007   0.0001  10.3846 0.0000   0.0006   0.00
08
tenure                 -0.0938  0.0063 -14.8132 0.0000  -0.1062  -0.08
14
MonthlyCharges         -0.0181  0.0020  -9.0806 0.0000  -0.0220  -0.01
42
OnlineSecurity_No      0.4323   0.0927   4.6616 0.0000   0.2505   0.61
40
PaymentMethod_Electronic check 0.4507   0.0809   5.5725 0.0000   0.2922   0.60
92
TechSupport_No         0.3935   0.0920   4.2770 0.0000   0.2132   0.57
39
InternetService_Fiber optic 1.2233   0.1312   9.3271 0.0000   0.9662   1.48
04
=====
==
80.88026502602933%
```

**We will now compare two other popular methods: Decision Tree, and K-Nearest-Neighbor to gauge best model to predict**

```
In [272]: #Print weights ov variables
weights = pd.Series(model.coef_[0], index = X_feat.columns.values)
weights.sort_values(ascending = False)
weights.sort_values(ascending = True)
```

```
Out[272]: Contract_Two year      -0.696976
tenure      -0.063308
TotalCharges      0.000307
MonthlyCharges      0.000496
SeniorCitizen      0.329954
PaymentMethod_Electronic check      0.427774
OnlineSecurity_No      0.433929
TechSupport_No      0.436380
InternetService_Fiber optic      0.897562
dtype: float64
```

## Exploring other Machine Learning Predictive Models

### Using decision trees then k-nearest-neight

```
In [304]: from sklearn import tree

model = tree.DecisionTreeClassifier(random_state=0, max_depth=5)

from sklearn import metrics

results = []
for i in range(1,10): # replicating model to get best average
    result = model.fit(X_train, Y_train)
    prediction_test = model.predict(X_test)
    #print accuracy
    results.append(float(metrics.accuracy_score(Y_test, prediction_test)*100))

#print((results))
#another method to get accuracy:
#from sklearn.metrics import accuracy_score
print(str(sum((results))/len(results))+ '%')
#accuracy_score(Y_test, prediction_test)

78.98722195929957%
```



Not as good as our logistics regression model, but did not hurt to check

```
In [275]: from sklearn.neighbors import KNeighborsClassifier

score_knn = 0
neighborcount = 3
max_value = -1

while score_knn > max_value:
    model = KNeighborsClassifier(n_neighbors=neighborcount)
    result = model.fit(X_train, Y_train)
    prediction_test = model.predict(X_test)
    max_value = float(score_knn)
    score_knn = float(metrics.accuracy_score(Y_test, prediction_test))
    neighborcount +=1
print(str(max_value*100)+'% using '+str(neighborcount) + ' neighbors')
#print((results))
#another method to get accuracy:
#from sklearn.metrics import accuracy_score

#accuracy_score(Y_test, prediction_test)
```

76.99952673923332% using 8 neighbors

A little better than our decision tree model, but still not as good as our linear regression.

## End of Case Study

**[Personal Website \(https://technerdhelp.com\)](https://technerdhelp.com)**

In [ ]: